

Prunable Authenticated Log and Authenticable Snapshot in Distributed Collaborative Systems


Victorien Elvinger

October 2018



Supervised by G rald Oster and Fran ois Charoy


Peer-to-Peer Collaborative Systems



- participants issue **operations** to update a shared document

Peer-to-Peer Collaborative Systems


2



- participants issue **operations** to update a shared document


Peer-to-Peer Collaborative Systems

3



- participants issue **operations** to update a shared document
- participants have their **own copy** of **co-authored data**


Peer-to-Peer Collaborative Systems



- participants issue **operations** to update a shared document
- participants have their **own copy** of **co-authored data**
 - **always-available**
 - continuously **diverge then converge** from others


Peer-to-Peer Collaborative Systems

5




- participants issue **operations** to update a shared document
- participants have their **own copy** of **co-authored data**
 - **always-available**
 - continuously **diverge then converge** from others

Peer-to-Peer Collaborative Systems



- participants issue **operations** to update a shared document
- participants have their **own copy** of **co-authored data**
 - **always-available**
 - continuously **diverge then converge** from others
 - convergence as a **liveness property**


Threat to convergence



- Byzantine^[1] adversary
 - **colluding malicious participants + network**
 - can **conceal divergences** using **equivocation**
- unlimited exchanges between honest participants

^[1]Leslie Lamport et al. The Byzantine Generals Problem. *ACM Transactions on Programming Languages and Systems (TOPLAS)*, July 1982.


Threat to convergence



- Byzantine^[1] adversary
 - **colluding malicious participants + network**
 - can **conceal divergences** using **equivocation**
- unlimited exchanges between honest participants

^[1]Leslie Lamport et al. The Byzantine Generals Problem. *ACM Transactions on Programming Languages and Systems (TOPLAS)*, July 1982.


Threat to convergence



- Byzantine^[1] adversary
 - **colluding malicious participants + network**
 - can **conceal divergences** using **equivocation**
- unlimited exchanges between honest participants

^[1]Leslie Lamport et al. The Byzantine Generals Problem. *ACM Transactions on Programming Languages and Systems (TOPLAS)*, July 1982.


Threat to convergence



- Byzantine^[1] adversary
 - **colluding malicious participants + network**
 - can **conceal divergences** using **equivocation**
- unlimited exchanges between honest participants

^[1]Leslie Lamport et al. The Byzantine Generals Problem. *ACM Transactions on Programming Languages and Systems (TOPLAS)*, July 1982.


Threat to convergence



- Byzantine^[1] adversary
 - **colluding malicious participants + network**
 - can **conceal divergences** using **equivocation**
- unlimited exchanges between honest participants

^[1]Leslie Lamport et al. The Byzantine Generals Problem. *ACM Transactions on Programming Languages and Systems (TOPLAS)*, July 1982.


Authenticated logs



- **misbehaviors are accountable**


- log entries are signed by their author
- secured dependency tracking of operations thanks to hashes

Authenticated logs




- **misbehaviors are accountable**
 - log entries are signed by their author
 - secured dependency tracking of operations thanks to hashes
- **strong convergence**
 - convergent logs \implies convergent documents

Overheads of authenticated logs




- participants **keep the entire log** for newcomers

Overheads of authenticated logs




- participants **keep the entire log** for newcomers

Overheads of authenticated logs



- participants **keep the entire log** for newcomers
- **newcomers** verify and play the **entire log**


Overheads of authenticated logs



- participants **keep the entire log** for newcomers
- **newcomers** verify and play the **entire log**


**How to reduce overheads of authenticated
logs in dynamic groups?**

Overview of contributions




- participants **prune their log**

Overview of contributions




- participants **prune their log**

Overview of contributions



- participants **prune their log**
- newcomers **retrieve a document and its pruned log**
 1. verify the **consistency** of the pruned log
 2. authenticate the document using the pruned log


Overview of contributions



- participants **prune their log**
- newcomers **retrieve a document and its pruned log**
 1. verify the **consistency** of the pruned log
 2. authenticate the document using the pruned log

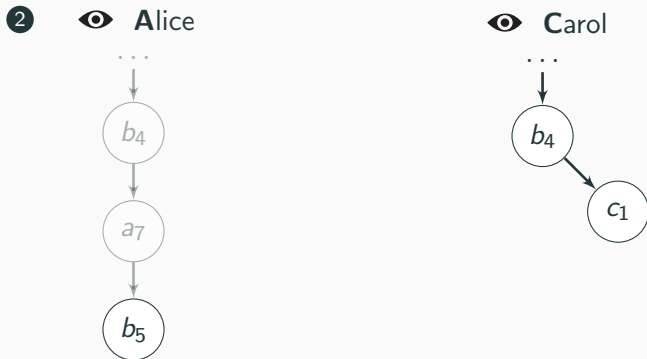
**How to prune a log without threatening
convergence?**

Pruning issue



- operations **cannot be arbitrarily dropped**
 - operations are required to verify a given operation
- concurrency plays an important role

Pruning issue




- operations **cannot be arbitrarily dropped**
 - operations are required to verify a given operation
- concurrency plays an important role

Definition


An operation is stable once it is no longer possible to accept a concurrent operation to it.

- primitive to define which operations can be dropped
 - stable part can be pruned
- necessity of **limiting concurrencies**

Consistency models and Causal consistency




Causal consistent log



Non causal consistent log

- logs respect a **consistency model**
 - how operations can be related
- Causal consistency
 - Participants **linearly order** their operations

👁 Bob group = {A, B}




Definition

An operation is *causally stable* (cs) once all participants of the group have observed it.

^[2]Carlos Baquero et al. Making Operation-Based CRDTs Operation-Based. In *Distributed Applications and Interoperable Systems - 14th IFIP WG 6.1 International Conference, DAIS 2014, Held as Part of the 9th International Federated Conference on Distributed Computing Techniques*. Springer, 2014 .

Limitations of Causal Stability


1 👁 Bob group = $\{A, B\}$



Issues

Limitations of Causal Stability

2 👁 **Alice** group = {A, B}




Issues

- convergent logs do not imply **convergent stability**

Limitations of Causal Stability


3 👁 Bob group = ?



Issues


- convergent logs do not imply **convergent stability**
- does not fit **dynamic groups**

Membership tracking



- Operations embed a list of the group members
- Dependencies of an operation must include its author

Contribution: Provable causal stability




Definition

An operation is provably causally stable (pcs) once all required observers have provably observed it.

required observers =

- + participants at the moment of generation
- + participants concurrently invited

Contribution: Provable causal stability




Definition

An operation is provably causally stable (pcs) once all required observers have provably observed it.

required observers =

- + participants at the moment of generation
- + participants concurrently invited

Contribution: Provable causal stability




Definition

An operation is provably causally stable (pcs) once all required observers have provably observed it.

required observers =

- + participants at the moment of generation
- + participants concurrently invited

Contribution: Provable causal stability




Definition

An operation is provably causally stable (pcs) once all required observers have provably observed it.

required observers =


- + participants at the moment of generation
- + participants concurrently invited

Malicious participants and causal consistency




- Malicious participants can issue **non-linear operations**
 - Causal consistency is not achievable

View-Fork-Join-Causal consistency^[3]



VFJC-consistent



VFJC-consistent

- honest participants linearly order their operations
- operations **cannot directly depend on non-linear** ones


^[3]Prince Mahajan et al. Consistency, Availability, and Convergence. Technical report TR-11-22, The University of Texas at Austin, 2011.

View-Fork-Join-Causal consistency^[3]

2




Alice



non VFJC-consistent



Bob



VFJC-consistent

- honest participants linearly order their operations
- operations **cannot directly depend on non-linear** ones


^[3]Prince Mahajan et al. Consistency, Availability, and Convergence. Technical report TR-11-22, The University of Texas at Austin, 2011.

View-Fork-Join-Causal consistency^[3]

3




Alice



VFJC-consistent



Bob




VFJC-consistent

- honest participants linearly order their operations
- operations **cannot directly depend on non-linear** ones


^[3]Prince Mahajan et al. Consistency, Availability, and Convergence. Technical report TR-11-22, The University of Texas at Austin, 2011.

VFJC consistency



- **impossibility of** honestly accepting concurrent operations which are **non-linear with a provably causally stable** operation

Contribution: VFJC Stability



Definition

An operation is VFJC-stable (vfjcs) once all required observers have a provably causally stable operation which depends on.

- VFJC stability \implies provable causal stability

Stability

Stability	CS	PCS	VFJCS	VFJCS
Consistency model	Causal	\sim Causal	VFJC	SVFJC
Group membership	Static	Dynamic	Static	Dynamic
environment	Honest	Honest	Malicious	Malicious

CS: Causal Stability

PCS: Provable Causal Stability

VFJC: View-Fork-Join-Causal

VFJCS: View-Fork-Join-Causal Stability

SVFJC: Stabilizable View-Fork-Join-Causal

Contributions

- adaptation of **Causal Stability** for dynamic groups
- definition of **VFJC Stability**
- Stabilizable View-Fork-Join-Causal (VFJC) (not covered here)
- document **authentication based on a pruned log**

Perspectives

- exploring sub-group log pruning
- VFJC Stability as building block of **asynchronous Byzantine consensus**


References

- [1] Carlos Baquero et al. Making Operation-Based CRDTs Operation-Based. In *Distributed Applications and Interoperable Systems - 14th IFIP WG 6.1 International Conference, DAIS 2014, Held as Part of the 9th International Federated Conference on Distributed Computing Techniques*. Springer, 2014.
- [2] Prince Mahajan et al. Consistency, Availability, and Convergence. Technical report TR-11-22, The University of Texas at Austin, 2011.
- [3] Leslie Lamport et al. The Byzantine Generals Problem. *ACM Transactions on Programming Languages and Systems (TOPLAS)*, July 1982.



How to authenticate an untrusted snapshot?

Overview of document authentication



- participants embed a **document fingerprint** in invitations
 - participants verify fingerprints
- **computed** fingerprint = **embedded** fingerprint
⇒ **document considered authentic**
- a newcomer trusts a **fingerprint verified by her trusted set**